

# BUSINESS IMPACT

[SIGN UP TO RECEIVE BLOG UPDATES](#)[SORT POSTS BY TOPIC](#)**JUN 18, 2019**

## PREDICTING SALES WITH THE AID OF PANDAS

POSTED BY [MEGAN QUINN](#)

[Pandas](#) is an open-source Python package that provides users with high-performing and flexible data structures. These structures are designed to make analyzing relational or labeled data both easy and intuitive. Pandas is one of the most popular and quintessential tools leveraged by data scientists when developing a machine learning model. The most crucial step in the machine learning process is not simply fitting a model to a given data set. Most of the model development process takes place in the pre-processing and data exploration phase. An accurate model requires good predictors and, in order to acquire them, the user must understand the raw data. Through Pandas' numerous data wrangling and analysis tools, this important step can easily be achieved. The goal of this blog is to highlight some of the central and most commonly used tools in Pandas while illustrating their significance in model development. The data set used for this demo consists of a supermarket chain's sales across multiple stores in a variety of cities. The sales data is broken down by items within the stores. The goal is to predict a certain item's sale.

### Reading the Data

When starting a new Python script, modules required for the analysis, including Pandas, must be imported into the environment:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

### Latest Posts

[Predicting Sales with the Aid of Pandas](#)[Consider 4 Factors to Amplify Health Analytics Efforts](#)[Meet Our Team: Leo Furlong](#)[Knowledge Mining Showcase: Azure Search](#)[BlueGranite Joins Developers and Industry Experts at Microsoft Build 2019](#)[BlueGranite is a Great Place to Work in 2019 - Here's Why!](#)[Create a Dynamic Title in Power BI \(Updated\)](#)

FREE WHITEPAPER

**Four Factors To Consider Before Implementing AI**[LEARN MORE](#)

in the current working directory folder for the Python environment.

```
StoreSales_df=pd.read_csv('StoreSales.csv')
```

Once the data frame is created, there are a variety of viewing and inspecting tools available in order to achieve a better understanding of the raw data. The `df.head(n)` and `df.tail(n)` functions allow users to examine the first and last rows respectively:

```
StoreSales_df.head(5)
StoreSales_df.tail(5)
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Size	Outlet_Location
0	FDA15	9.30	Low Fat	0.016047	Dairy	249.8092	OUT049	1999	Medium	
1	DRC01	5.92	Regular	0.019278	Soft Drinks	48.2692	OUT018	2009	Medium	
2	FDN15	17.50	Low Fat	0.016760	Meat	141.6180	OUT049	1999	Medium	
3	FDX07	19.20	Regular	0.000000	Fruits and Vegetables	182.0950	OUT010	1998	NaN	
4	NCD19	8.93	Low Fat	0.000000	Household	53.8614	OUT013	1987	High	

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Establishment_Year	Outlet_Size	Outlet_Location
8518	FDF22	6.865	Low Fat	0.056783	Snack Foods	214.5218	OUT013	1987	High	
8519	FDS36	8.380	Regular	0.046982	Baking Goods	108.1570	OUT045	2002	NaN	
8520	NCJ29	10.600	Low Fat	0.035186	Health and Hygiene	85.1224	OUT035	2004	Small	
8521	FDN46	7.210	Regular	0.145221	Snack Foods	103.1332	OUT018	2009	Medium	
8522	DRG01	14.800	Low Fat	0.044878	Soft Drinks	75.4670	OUT046	1997	Small	

The `df.shape` and `df.info` provide information about the number of rows and columns in a data frame, the data types, and missing data:

```
StoreSales_df.shape
```

```
(8523, 12)
```

```
StoreSales_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8523 entries, 0 to 8522
Data columns (total 12 columns):
Item_Identifier      8523 non-null object
Item_Weight          7060 non-null float64
Item_Fat_Content      8523 non-null object
Item_Visibility       8523 non-null float64
Item_Type            8523 non-null object
Item_MRP             8523 non-null float64
Outlet_Identifier     8523 non-null object
Outlet_Establishment_Year 8523 non-null int64
Outlet_Size          6113 non-null object
Outlet_Location_Type  8523 non-null object
Outlet_Type          8523 non-null object
Item_Outlet_Sales     8523 non-null float64
dtypes: float64(4), int64(1), object(7)
memory usage: 799.1+ KB
```

```
min(StoreSales_df.Outlet_Establishment_Year)
max(StoreSales_df.Outlet_Establishment_Year)
```

## Data Cleaning and Feature Engineering

After getting a general overview and understanding of the data, the next step toward successful model development is cleaning the data and creating new, possibly more influential variables from the existing raw data.

The variable *Item\_Identifier* follows a labeling pattern of letters per each product (i.e., 'FD' for food, 'DR' for drinks and 'NC' for n) followed by a three-digit code. It may be more useful to have a group type variable with just these two letters rather than the entire code. To achieve this, the `.map()` function applies a selection of only the first two values in the item identifier and returns as a new column labeled *Item\_Group\_Type*.

```
In [12]: StoreSales_df['Item_Group_Type']=StoreSales_df.Item_Identifier.map(lambda x: x[:2])
StoreSales_df[['Item_Identifier','Item_Group_Type']].head(5)
```

```
Out[12]:
```

	Item_Identifier	Item_Group_Type
0	FDA15	FD
1	DRC01	DR
2	FDN15	FD
3	FDX07	FD
4	NCD19	NC

In order to analyze the outlet establishment year as a numerical variable, a new column entitled *Outlet\_Age* can be calculated by subtracting the outlet's year by the the max year value of the dataset plus one (assuming this data's collection ended the prior year). The *Outlet\_Establishment\_Year* variable can then be dropped using the Pandas `df.drop()` function.

```
StoreSales_df["Outlet_Age"]=(max(StoreSales_df.Outlet_Establishment_Year)+1)-StoreSales_df["Outlet_Establishment_Year"]
StoreSales_df=StoreSales_df.drop(columns="Outlet_Establishment_Year")
StoreSales_df.head()
```

```
Out[13]:
```

Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Size	Outlet_Location_Type	Outlet_Type	Item_Outlet_Sales	Item_Group_Type	Outlet_Age
Low Fat	0.016047	Dairy	249.8092	OUT049	Medium	Tier 1	Supermarket Type1	3735.1380	FD	11
Regular	0.019278	Soft Drinks	48.2692	OUT018	Medium	Tier 3	Supermarket Type2	443.4228	DR	1
Low Fat	0.016760	Meat	141.6180	OUT049	Medium	Tier 1	Supermarket Type1	2097.2700	FD	11
Regular	0.000000	Fruits and Vegetables	182.0950	OUT010	NaN	Tier 3	Grocery Store	732.3800	FD	12
Low Fat	0.000000	Household	53.8614	OUT013	High	Tier 3	Supermarket Type1	994.7052	NC	23

Since most machine learning models in Python will return errors if null values exist within the data, identifying their existence and rectifying the issue is a crucial step. The code below counts the number of null values for each column in the store sales data frame:

```
Item_Fat_Content    0
Item_Visibility     0
Item_Type           0
Item_MRP            0
Outlet_Identifier    0
Outlet_Size        2410
Outlet_Location_Type 0
Outlet_Type         0
Item_Outlet_Sales    0
Item_Group_Type     0
Outlet_Age          0
dtype: int64
```

There are a variety of methods to address null values within a data set. Some common approaches include simply removing the rows containing the null values, forward or backward filling of values for timeseries data, or replacing the null value with a calculated value. This data set contains a significant number of null values. Removing the entire subject could lead to a lack of complete data, therefore filling in the missing values is a more appropriate method.

Since the *Item\_Weight* variable is numeric, replacing a null value with the item's average weight is a logical approach. The first line of the code uses Pandas `df.groupby()` function combined with the `.agg` function to find the mean weight of each unique item and store the results in another Pandas data frame. In lines two and three, setting the index of the original data frame to match the index of the new *Item\_Identifier\_Mean* data frame allows the null values to be easily imputed with their matching mean values. It is then necessary to check to see if this method resolved all the null values. Line 5 of the code reveals that four rows still contain null values.

```
In [8]: Item_Identifier_Mean=StoreSales_df.groupby("Item_Identifier").agg({'Item_Weight': 'mean'})
```

```
StoreSales_df2=StoreSales_df.set_index("Item_Identifier")
StoreSales_df2.Item_Weight.fillna(Item_Identifier_Mean.Item_Weight, inplace= True)
StoreSales_df2=StoreSales_df2.reset_index()

StoreSales_df2[pd.isnull(StoreSales_df2.Item_Weight)]
```

```
Out[8]:
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Size	Outlet_Location_Type	Outlet_Type	Item
927	FDN52	NaN	Regular	0.130933	Frozen Foods	86.9198	OUT027	Medium	Tier 3	Supermarket Type3	
1922	FDK57	NaN	Low Fat	0.079904	Snack Foods	120.0440	OUT027	Medium	Tier 3	Supermarket Type3	
4187	FDE52	NaN	Regular	0.029742	Dairy	88.9514	OUT027	Medium	Tier 3	Supermarket Type3	
5022	FDQ60	NaN	Regular	0.191501	Baking Goods	121.2098	OUT019	Small	Tier 1	Grocery Store	

To investigate the cause of this persisting issue, the new mean imputed data frame is merged on *Item\_Identifier* with the original data frame using Pandas `df.merge()`. The merge reveals that the items with null values in the new data frame only appeared once in the original data and had no weight information, therefore a mean could not be calculated.

2	FDE52	NaN	Regular	0.029742	Dairy	88.9514	OUT027	Medium	Tier 3
3	FDQ60	NaN	Regular	0.191501	Baking Goods	121.2098	OUT019	Small	Tier 1

4 rows × 25 columns

This same method is then applied to the *Outlet\_Size* variable, except the mode of the outlet type is used as the imputed value since *Outlet\_Size* is categorical. To calculate the mode, the outlet types are converted into a Pandas structure called a series, then the mode is applied to each of these series using the `.apply` function. From there, a similar merge as before can take place. However, since the merge is on a column name instead of the index, both *Outlet\_Size* columns are retained in the new dataset with “\_x” and “\_y” appended to the names. Only one of these columns is needed, therefore renaming to remove the “\_x” and dropping the *Outlet\_Size\_y* column is conducted in line 4.

```
In [10]: Outlet_Identifier_Mode=pd.DataFrame(StoreSales_df2.groupby("Outlet_Type")['Outlet_Size'].apply(pd.Series.mode))
StoreSales_df3=StoreSales_df2.merge(Outlet_Identifier_Mode,on="Outlet_Type")
StoreSales_df3['Outlet_Size_x'].fillna(StoreSales_df3.Outlet_Size_y,inplace=True)
StoreSales_df3=StoreSales_df3.rename(columns={'Outlet_Size_x': 'Outlet_Size'}).drop(columns=['Outlet_Size_y'])
StoreSales_df3
```

Out[10]:

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_MRP	Outlet_Identifier	Outlet_Size	Outlet_Location_Type	Outlet_Type
0	FDA15	9.300	Low Fat	0.016047	Dairy	249.8092	OUT049	Medium	Tier 1	Supermarket Type1
1	FDN15	17.500	Low Fat	0.016760	Meat	141.6180	OUT049	Medium	Tier 1	Supermarket Type1
2	NCD19	8.930	Low Fat	0.000000	Household	53.8614	OUT013	High	Tier 3	Supermarket Type1
3	FDO10	13.650	Regular	0.012741	Snack Foods	57.6588	OUT013	High	Tier 3	Supermarket Type1
4	FDH17	16.200	Regular	0.016687	Frozen Foods	96.9726	OUT045	Small	Tier 2	Supermarket Type1
5	FDU28	19.200	Regular	0.094450	Frozen Foods	187.8214	OUT017	Small	Tier 2	Supermarket Type1
6	FDY07	11.800	Low Fat	0.000000	Fruits and Vegetables	45.5402	OUT049	Medium	Tier 1	Supermarket Type1

Examining the new cleaned data shows that *Outlet\_Type* no longer contains null values.

```
In [11]: pd.isnull(StoreSales_df3.Outlet_Type).sum()
```

Out[11]: 0

Returning to the null *Item\_Weight* values, since there were four items with only one record each, these values can simply be dropped. Checking the final cleaned data set reveals that all null values have been corrected:

```

dtype: object
Item_MRP      0
Outlet_Identifier  0
Outlet_Size    0
Outlet_Location_Type  0
Outlet_Type    0
Item_Outlet_Sales  0
Item_Group_Type  0
Outlet_Age     0
dtype: int64

```

Dividing the features into categorical and numerical sets allows for further examination of any other possible incongruities in the data. Using the `df.select_dtypes()` function, two data frames are created containing only the specified data types:

```

In [13]: categorical_features = StoreSales_Cleaned.select_dtypes(include=[np.object])
numerical_features = StoreSales_Cleaned.select_dtypes(include=[np.number])

categorical_features.dtypes
numerical_features.dtypes

```

```

Out[13]: Item_Identifier      object
Item_Fat_Content            object
Item_Type                   object
Outlet_Identifier           object
Outlet_Size                 object
Outlet_Location_Type        object
Outlet_Type                 object
Item_Group_Type             object
dtype: object

```

```

Out[13]: Item_Weight          float64
Item_Visibility              float64
Item_MRP                     float64
Item_Outlet_Sales            float64
Outlet_Age                   int64
dtype: object

```

Through the use of a for loop and the `.unique()` function, the number of unique values for each categorical feature and the label can be displayed. As is fairly common with categorical variables recorded from different sources, the labeling technique of *Item\_Fat\_Content* is inconsistent. “Low Fat” is represented as both “low fat” and “LF” while “Regular” is also recorded as “reg”. This type of discrepancy will cause issues when creating dummy variables.

```

In [14]: for name in list(categorical_features):
print (name + ":")
print ("    Count of unique values:" + str(len(categorical_features[name].unique())))
print (categorical_features[name].unique())

```

```

Item_Identifier:
    Count of unique values:1555
['FDA15' 'FDN15' 'NCD19' ... 'FDC23' 'FDR07' 'FDP15']
Item_Fat_Content:
    Count of unique values:5
['Low Fat' 'Regular' 'low fat' 'reg' 'LF']
Item_Type:
    Count of unique values:16
['Dairy' 'Meat' 'Household' 'Snack Foods' 'Frozen Foods'
'Fruits and Vegetables' 'Breakfast' 'Hard Drinks' 'Breads' 'Soft Drinks'
'Health and Hygiene' 'Canned' 'Baking Goods' 'Starchy Foods' 'Others'
'Seafood']
Outlet_Identifier:
    Count of unique values:10
['OUT049' 'OUT013' 'OUT045' 'OUT017' 'OUT046' 'OUT035' 'OUT018' 'OUT010'
'OUT019' 'OUT027']
Outlet_Size:
    Count of unique values:3
['Medium' 'High' 'Small']
Outlet_Location_Type:
    Count of unique values:3
['Tier 1' 'Tier 3' 'Tier 2']
Outlet_Type:
    Count of unique values:4
['Supermarket Type1' 'Supermarket Type2' 'Grocery Store'
'Supermarket Type3']
Item_Group_Type:
    Count of unique values:3
['FD' 'NC' 'DR']

```

```
Out[15]: array(['Low Fat', 'Regular'], dtype=object)
```


## Data Exploration

With the data cleaning and feature engineering completed, an even closer examination of the data and its relations can be conducted using Pandas in conjunction with the [Matplotlib](#) library. The histogram below shows that the target variable, *Item\_Outlet\_Sales*, is right skewed:

```
In [16]: plt.hist(StoresSales_Cleaned.Item_Outlet_Sales, bins=25)
plt.xlabel("Item_Outlet_Sales")
plt.ylabel("Number of Sales")
plt.title("Item_Outlet_Sales Distribution")

Out[16]: (array([1.414e+03, 1.267e+03, 1.166e+03, 1.058e+03, 8.210e+02, 7.550e+02,
5.440e+02, 4.050e+02, 3.350e+02, 2.150e+02, 1.880e+02, 1.280e+02,
8.000e+01, 5.600e+01, 3.100e+01, 1.500e+01, 1.100e+01, 1.200e+01,
6.000e+00, 4.000e+00, 1.000e+00, 1.000e+00, 0.000e+00, 1.000e+00,
1.000e+00]),
array([ 33.29, 555.436992, 1077.583984, 1599.730976,
2121.877968, 2644.02496, 3166.171952, 3688.318944,
4210.465936, 4732.612928, 5254.75992, 5776.906912,
6299.053904, 6821.200896, 7343.347888, 7865.49488,
8387.641872, 8909.788864, 9431.935856, 9954.082848,
10476.22984, 10998.376832, 11520.523824, 12042.670816,
12564.817808, 13086.9648 ]),
<a list of 25 Patch objects>)

Out[16]: Text(0.5,0,'Item_Outlet_Sales')
Out[16]: Text(0,0.5,'Number of Sales')
Out[16]: Text(0.5,1,'Item_Outlet_Sales Distribution')
```




Applying the `.corr()` function to the numerical features data frame created earlier, provides a symmetrical data frame of the variables' correlations to each other. Plotting a heatmap of this correlation data frame highlights that *Item\_MRP* has the strongest, positive correlation with the target variable while *Item\_Visibility* has a negative correlation.

	Item_MRP	Item_Outlet_Sales	Outlet_Age
Item_MRP	0.025975	-0.001155	1.000000
Item_Outlet_Sales	0.013168	-0.126297	0.567803
Outlet_Age	0.013426	0.074325	-0.004599

```
In [18]: f, ax = plt.subplots(figsize=(12, 9))
sns.heatmap(corr, vmax=.8, square=True)

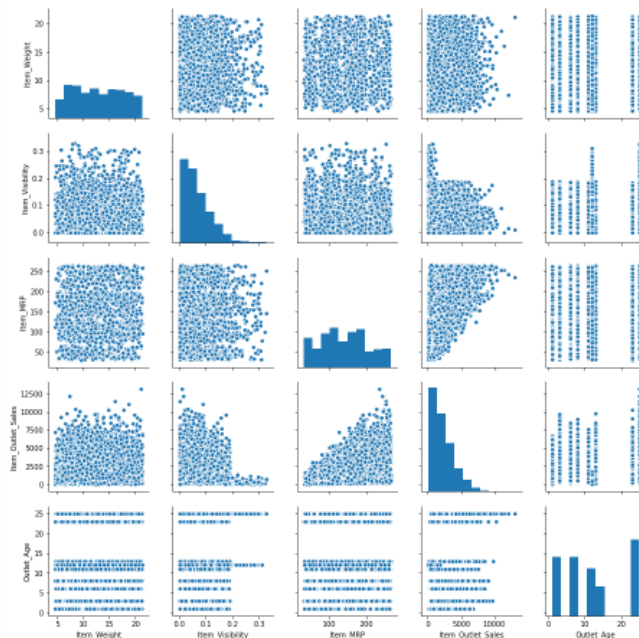
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x1fa501d2ba8>
```



Using the `.pairplot()` function allows for visualization of relationships among the all the numeric variables at once:

```
In [62]: sns.pairplot(StoreSales_Cleaned)

Out[62]: <seaborn.axisgrid.PairGrid at 0x1fa527c5c88>
```



The categorical variables' relationships with the target variable can also be examined through the use of `df.pivot_table()`. This function operates like pivot tables in Excel by creating an index and applying an aggregation function over a specified value. Plotting these



## Sales

```
In [19]: Item_Type_pivot = StoreSales_Cleaned.pivot_table(index='Item_Type', values='Item_Outlet_Sales', aggfunc=np.mean)
```

```
Item_Type_pivot.plot(kind='bar', color='blue', figsize=(25,7))
plt.xlabel("Item_Type")
plt.ylabel("Item_Outlet_Sales")
plt.title("Impact of Item_Type on Item_Outlet_Sales")
plt.xticks(rotation=0)
plt.show()
```


```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x1fa50296550>
```

```
Out[19]: Text(0.5,0,'Item_Type')
```

```
Out[19]: Text(0,0.5,'Item_Outlet_Sales')
```

```
Out[19]: Text(0.5,1,'Impact of Item_Type on Item_Outlet_Sales')
```

```
Out[19]: (array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15]),
<a list of 16 Text xticklabel objects>)
```



```
In [20]: Outlet_Size_pivot = StoreSales_Cleaned.pivot_table(index='Outlet_Size', values='Item_Outlet_Sales', aggfunc=np.mean)
```

```
Outlet_Size_pivot.plot(kind='bar', color='blue', figsize=(12,7))
plt.xlabel("Outlet_Size")
plt.ylabel("Item_Outlet_Sales")
plt.title("Impact of Outlet_Size on Item_Outlet_Sales")
plt.xticks(rotation=0)
plt.show()
```

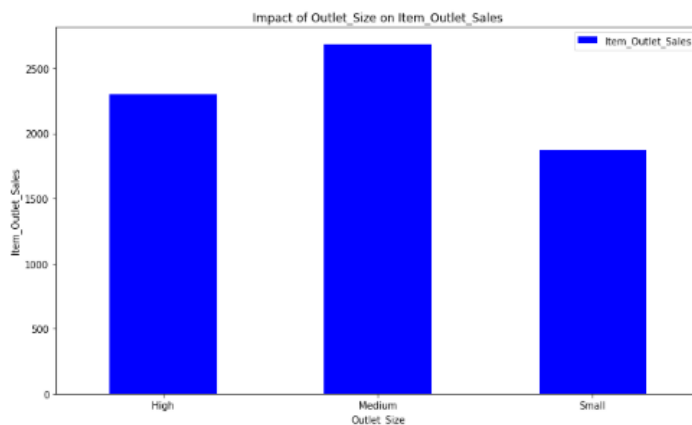
```
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x1fa501bfb48>
```

```
Out[20]: Text(0.5,0,'Outlet_Size')
```

```
Out[20]: Text(0,0.5,'Item_Outlet_Sales')
```

```
Out[20]: Text(0.5,1,'Impact of Outlet_Size on Item_Outlet_Sales')
```

```
Out[20]: (array([0, 1, 2]), <a list of 3 Text xticklabel objects>)
```



This data analysis and exploration will aid in guiding feature, model, and parameter selection during the model development phase.

Finally, the creation of dummy variables is required in order to use the categorical variables in the modeling process. Pandas has an easy to use function, `pd.get_dummies()`, that

```
In [22]: StoreSales_Dummy = pd.get_dummies(StoreSales_Cleaned, columns =['Item_Fat_Content','Outlet_Identifier','Outlet_Location_Type','Out
StoreSales_Dummy
```

n_Fat_Content_Low_Fat	Item_Fat_Content_Regular	Outlet_Identifier_OUT010	...	Outlet_Size_High	Outlet_Size_Medium	Outlet_Size_Small	Outlet_Type_Grocery_Store	Outlet_Type_Supermarket
1	0	0	...	0	1	0	0	0
1	0	0	...	0	1	0	0	0
1	0	0	...	1	0	0	0	0
0	1	0	...	1	0	0	0	0
0	1	0	...	0	0	1	0	0
0	1	0	...	0	0	1	0	0
1	0	0	...	0	1	0	0	0
0	1	0	...	0	0	1	0	0

## Model Development

With pre-processing, cleaning, and data exploration complete, the final phase of modeling can now take place. [Sklearn](#) is a commonly used machine learning library in Python that contains multiple modeling and evaluation tools. The first step is to enable the `train_test_split()` function of this package to divide the cleaned data frame into two separate data frames. The larger data frame, which will represent 85% of the entire data, will be used to train the model, while the remaining 15% will be used to evaluate and determine whether the model is appropriate.

```
In [26]: from sklearn.model_selection import train_test_split
train , test = train_test_split(StoreSales_Dummy,test_size=.15)
len(train)
len(test)
```

Out[26]: 7241

Out[26]: 1278

Next, the train and test data frames are each divided into two separate data frames, one containing the desired predictors and the other containing the target variable:

Out[45]:

	Item_Weight	Item_Visiblility	Item_MRP	Outlet_Age	Item_Fat_Content_Low Fat	Item_Fat_Content_Regular	Outlet_Identifier_OUT010	Outlet_Identifier_OUT013
7566	19.00	0.031024	210.5244	25	1	0	0	0
1236	9.80	0.047454	101.7016	3	1	0	0	0
7104	17.60	0.175546	163.6868	12	1	0	1	0
227	15.85	0.107765	59.5904	11	1	0	0	0
4200	15.30	0.022959	101.6332	23	1	0	0	1

5 rows x 29 columns

Out[45]:

	Item_Outlet_Sales
7566	1482.0708
1236	1518.0240
7104	163.7868
227	703.0848
4200	1845.5976

Since the target variable is continuous, a simple, yet standard approach is to test a linear regression model. Once imported from the sklearn package, the function is applied to the train data using the `model.fit()` function. The predictions are then stored in an array using `model.predict()`. Model evaluation is conducted by using a variety of the metric functions from sklearn, along with plotting the actual vs. predicted values. From these results, it appears linear regression may not be the best model for this data.

```
In [59]: from sklearn.linear_model import LinearRegression
```

```
lr = LinearRegression(normalize=True)
lr.fit(train_predictors, train_target)

train_predictions = lr.predict(train_predictors)
train_predictions
```

```
Out[59]: LinearRegression(copy_X=True, fit_intercept=True, n_jobs=None, normalize=True)
```

```
Out[59]: array([[1421.81051499],
 [1772.80318599],
 [ 663.53051985],
 ...,
 [2385.76875632],
 [2952.31715161],
 [2596.73167759]])
```

```
In [60]: from sklearn import metrics
```

```
np.sqrt(metrics.mean_squared_error(train_target, train_predictions))
metrics.r2_score(train_target, train_predictions)
metrics.mean_absolute_error(train_target, train_predictions)
```


```
Out[60]: 1120.252575629138
```

```
Out[60]: 0.5630719324461717
```

```
Out[60]: 831.0963748409376
```

```
In [61]: plt.scatter(train_target, train_predictions)
```

```
Out[61]: <matplotlib.collections.PathCollection at 0x1fa521644e0>
```



XGBRegressor tool is imported, it is applied to the train data and predictions are stored to evaluate the model. The metrics and plot both reveal higher performance than the linear regression model:

```
In [66]: from xgboost import XGBRegressor


xg_model = XGBRegressor(n_estimators=1000, learning_rate=0.05)
xg_model.fit(train_predictors, train_target)

xg_train_predictions = xg_model.predict(train_predictors)

Out[66]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
    colsample_bytree=1, gamma=0, importance_type='gain',
    learning_rate=0.05, max_delta_step=0, max_depth=3,
    min_child_weight=1, missing=None, n_estimators=1000, n_jobs=1,
    nthread=None, objective='reg:linear', random_state=0, reg_alpha=0,
    reg_lambda=1, scale_pos_weight=1, seed=None, silent=True,
    subsample=1)
```

```
In [71]: plt.scatter(train_target, xg_train_predictions)
```

```
Out[71]: <matplotlib.collections.PathCollection at 0x1fa68020198>
```



```
In [72]: np.sqrt(metrics.mean_squared_error(train_target, xg_train_predictions))
metrics.r2_score(train_target, xg_train_predictions)
metrics.mean_absolute_error(train_target, xg_train_predictions)
```

```
Out[72]: 922.0719026834847
```

```
Out[72]: 0.7039891577388857
```

```
Out[72]: 661.121754605931
```

This model then needs to be evaluated against the test data to determine if it is in fact a good model. The test data reveals that there is still some relevance to the model but further parameter tuning, and possibly other model selections may lead to better results.

```
In [73]: xg_test_predictions = xg_model.predict(test_predictors)


plt.scatter(test_target, xg_test_predictions)
np.sqrt(metrics.mean_squared_error(test_target, xg_test_predictions))
metrics.r2_score(test_target, xg_test_predictions)
metrics.mean_absolute_error(test_target, xg_test_predictions)
```

```
Out[73]: <matplotlib.collections.PathCollection at 0x1fa68075828>
```

```
Out[73]: 1141.6160734892937
```

```
Out[73]: 0.5843002719742592
```

```
Out[73]: 788.4657297851826
```



Overall, as evident by this demo, the actual fitting and tuning of a model is a small step compared to the entire machine learning process. In order to even obtain a data set for

also display the results of their analysis and modeling to their audience, making this library imperative for successful machine learning in Python.



FREE WHITEPAPER

## The Four Components To Consider When Implementing AI

LEARN MORE



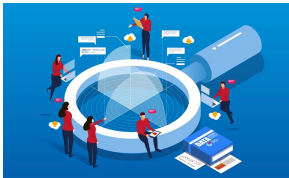


#### ABOUT THE AUTHOR


### MEGAN QUINN

Megan has expertise in statistical analysis and machine learning as well as statistical theory. Her recent focus has been centered on predictive maintenance for military fleets with a background in education research as well. She is knowledgeable in a variety of analytical tools including Python, R, SQL, and most recently Spark & Databricks.

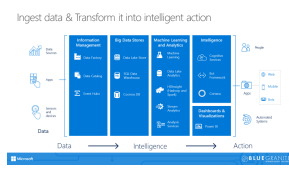
## RELATED ARTICLES



Knowledge Mining Showcase:  
Azure Search



Azure Data Lake Analytics Holds  
a Unique Spot in the Modern  
Data Architecture



Ingest data & Transform it into intelligent action

Cortana Intelligence Suite Tour

[ALL POSTS](#)[PREVIOUS POST](#)

Enterprise BI &  
Reporting  
Modern Data Platform

Events  
Resources

2750 Old Centre Rd Ste  
150  
Portage, MI 49024  
T. 877-817-0736

