

# Machine Learning and Big Data @ Uber: A Tale of Two Systems

Zhenxiao Luo

Engineering Manager @ Uber



# Agenda

Mission

Uber Business Highlights

Machine Learning @ Uber

Pain Points

Big Data @ Uber

Machine Learning support in Big Data

Ongoing Work



## Uber Mission

Transportation as reliable as running water, everywhere, for everyone



# Uber Stats

The Uber logo, consisting of the word "UBER" in white, uppercase, sans-serif font, centered within a solid black square.

UBER

6

Continents

73

Countries

450

Cities

12,000

Employees

10+ Million

Avg. Trips/Day

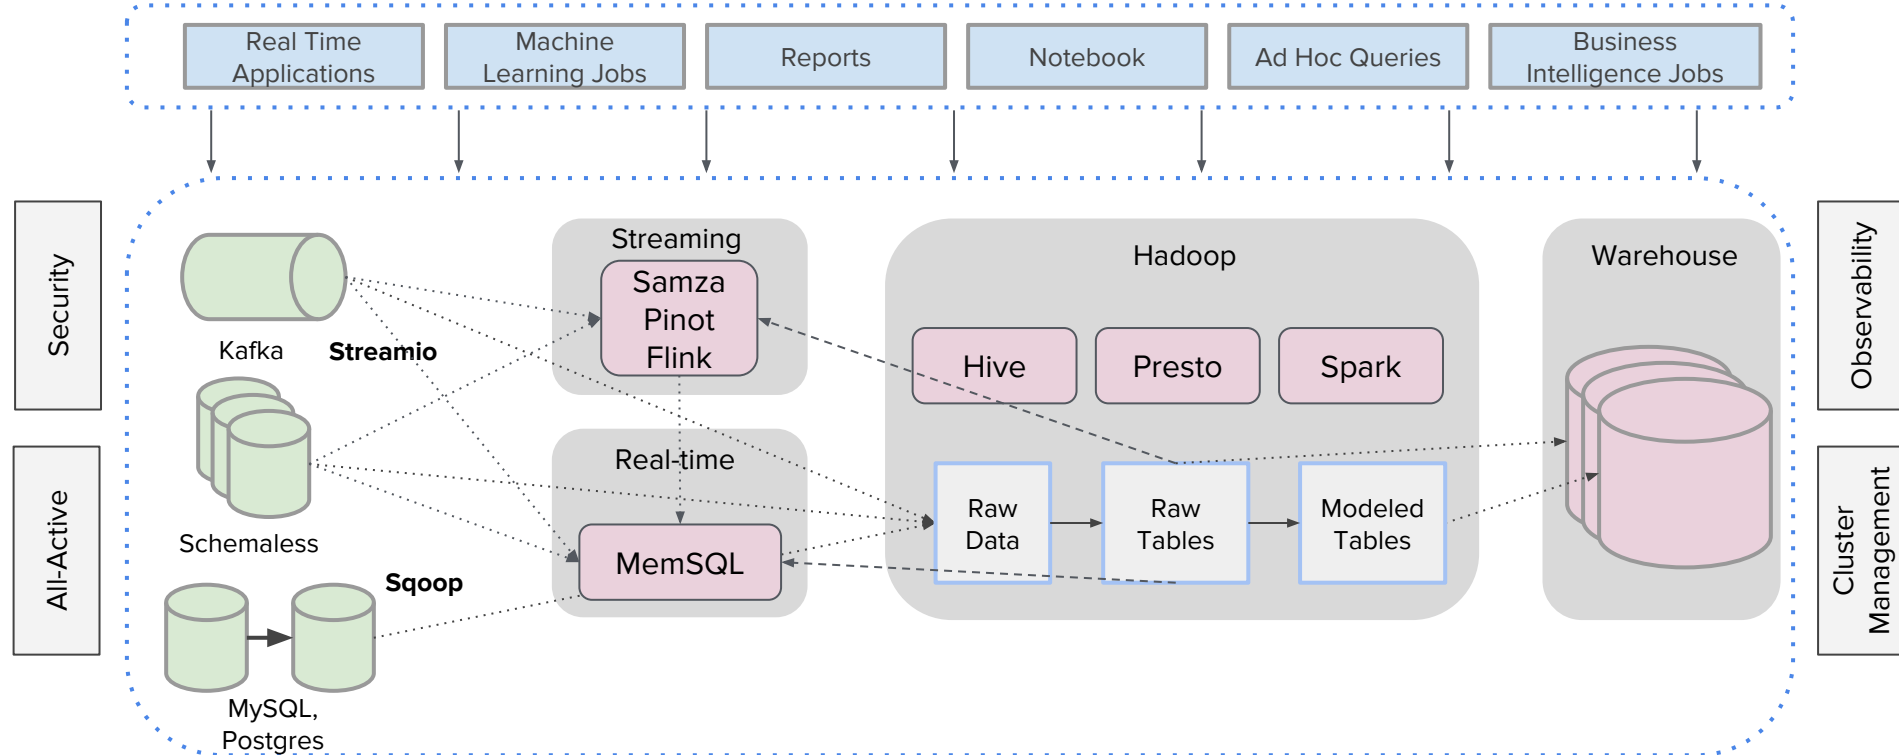
40+ Million

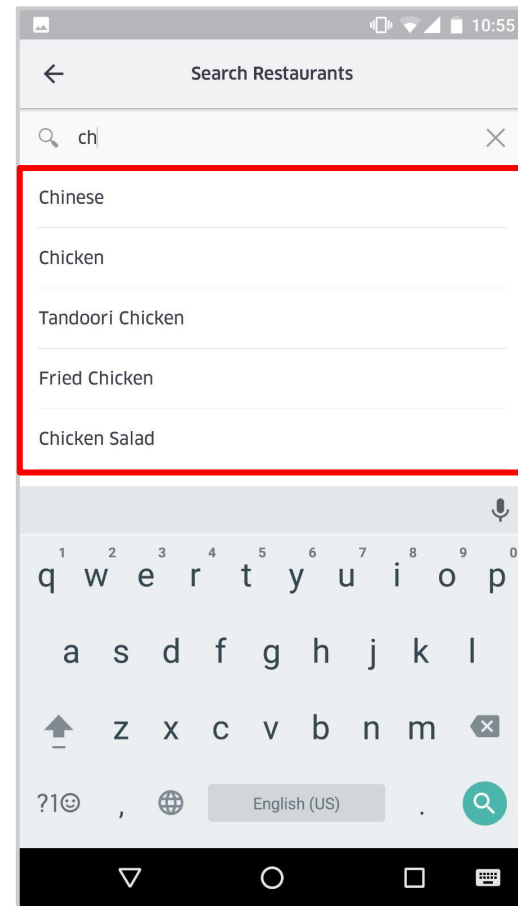
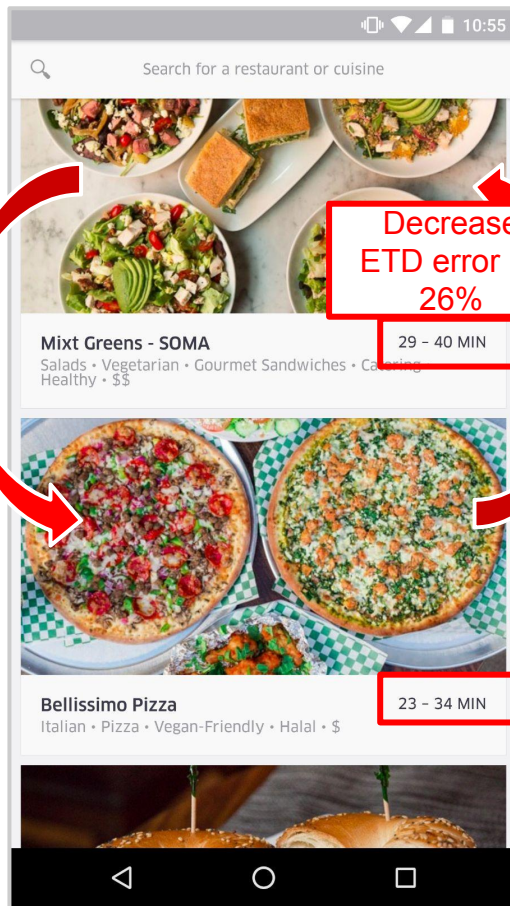
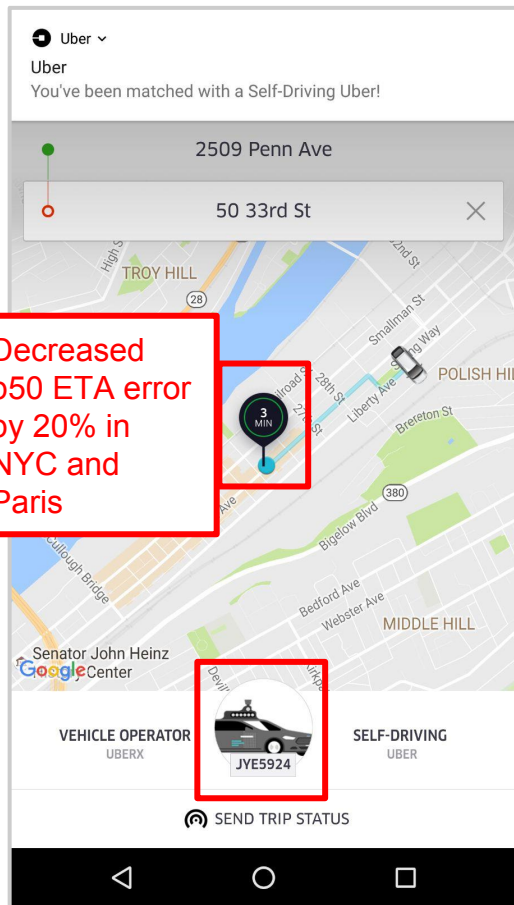
MAU Riders

1.5+ Million

MAU Drivers

# Data Infrastructure @ Uber



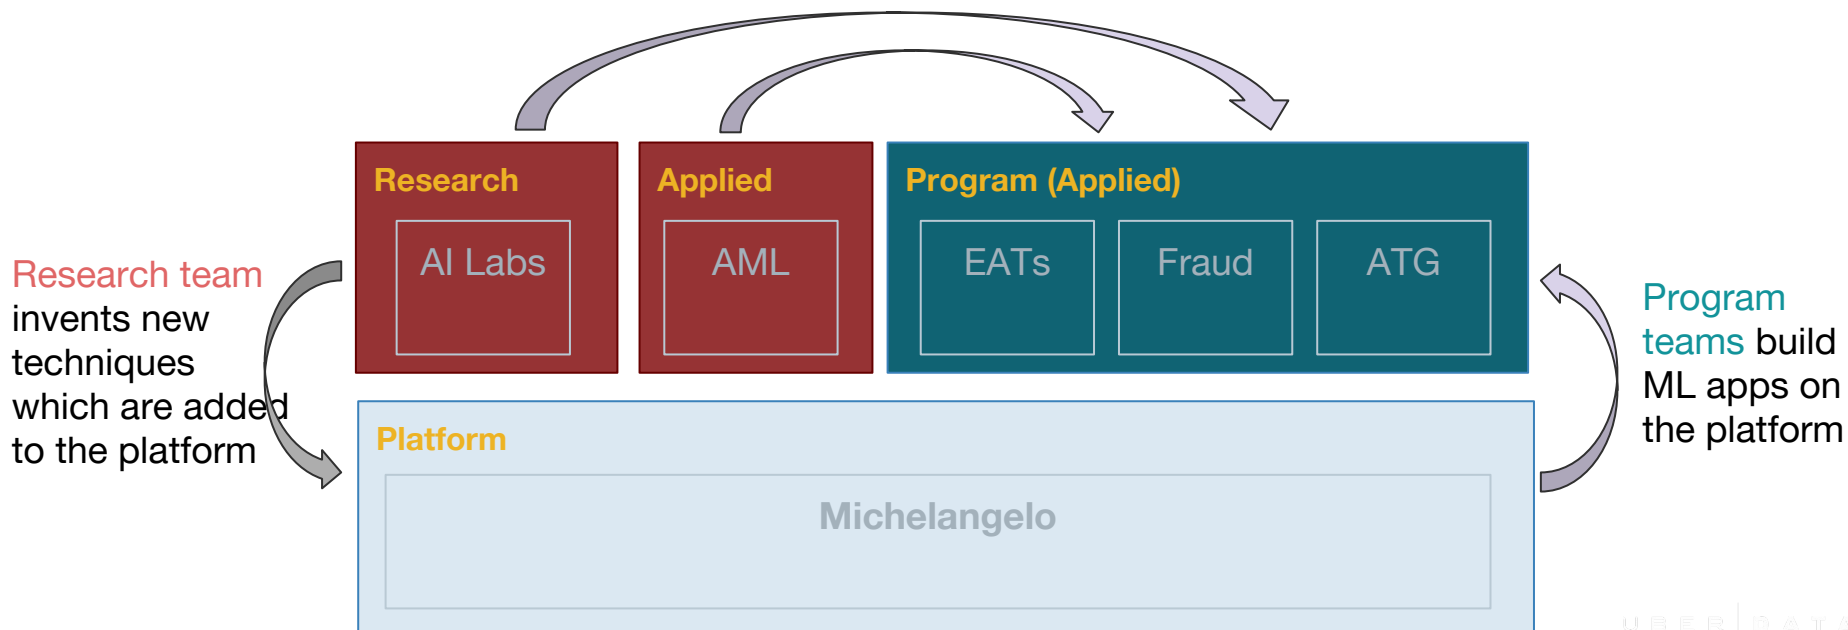




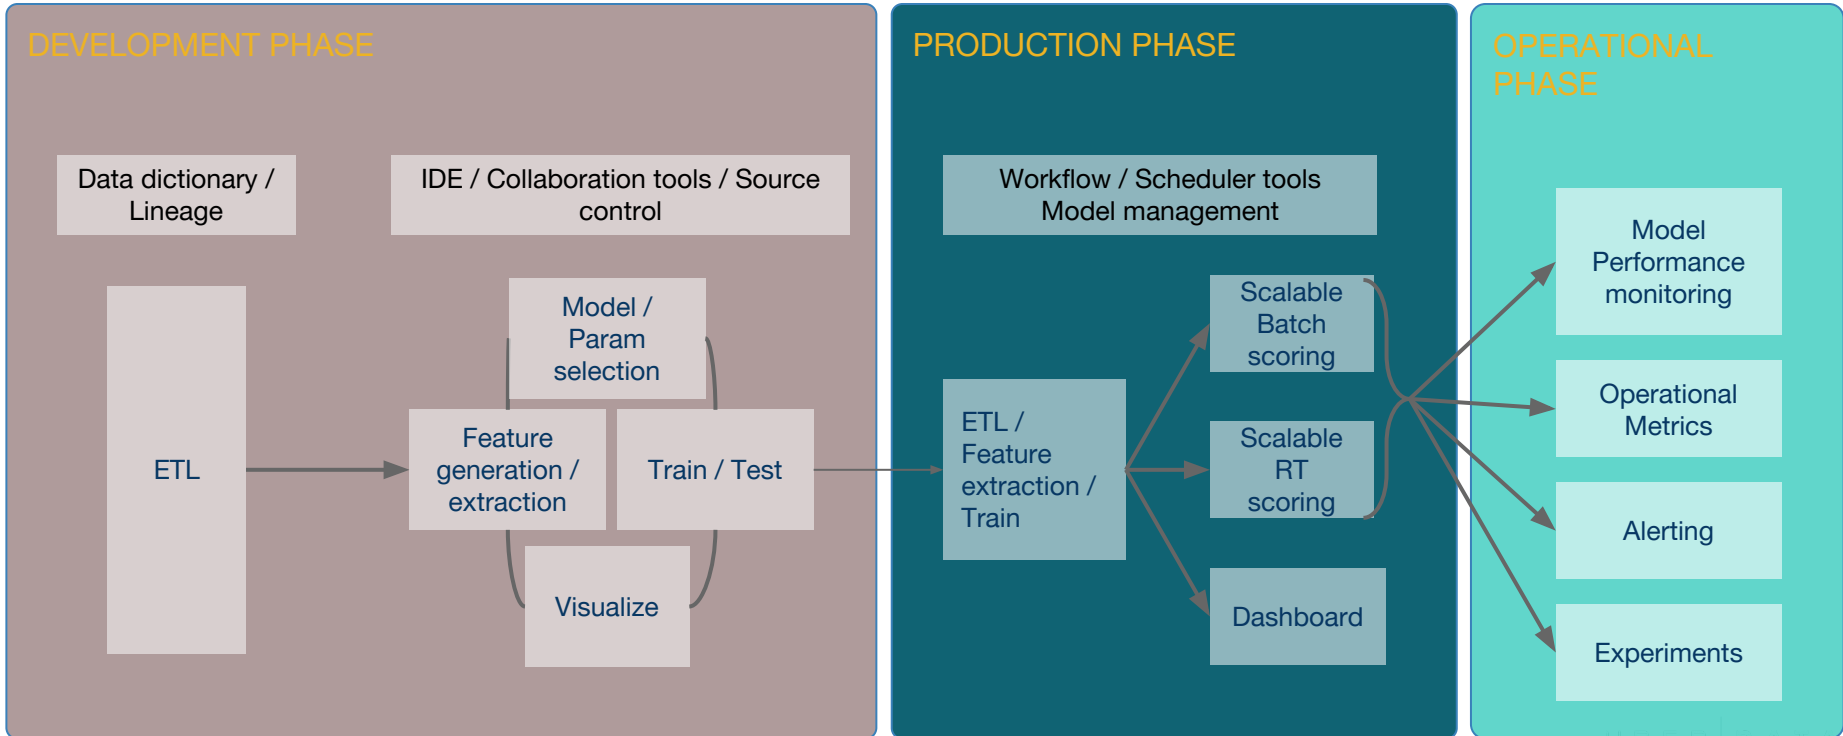
# Michelangelo: Uber Machine Learning Platform

**Research teams** focus their research based on business priorities from program teams

**Applied ML** teams helps with problems when Program team doesn't have right expertise; platform-izes the solutions



# Workflow of a Machine Learning Project






# Workflow

Same basic ML workflow & system requirements for

- Traditional ML & Deep Learning
- Supervised, Unsupervised & Semi-supervised learning
- Online learning
- Batch, Realtime and Mobile deployments
- Timeseries Forecasting



GET DATA

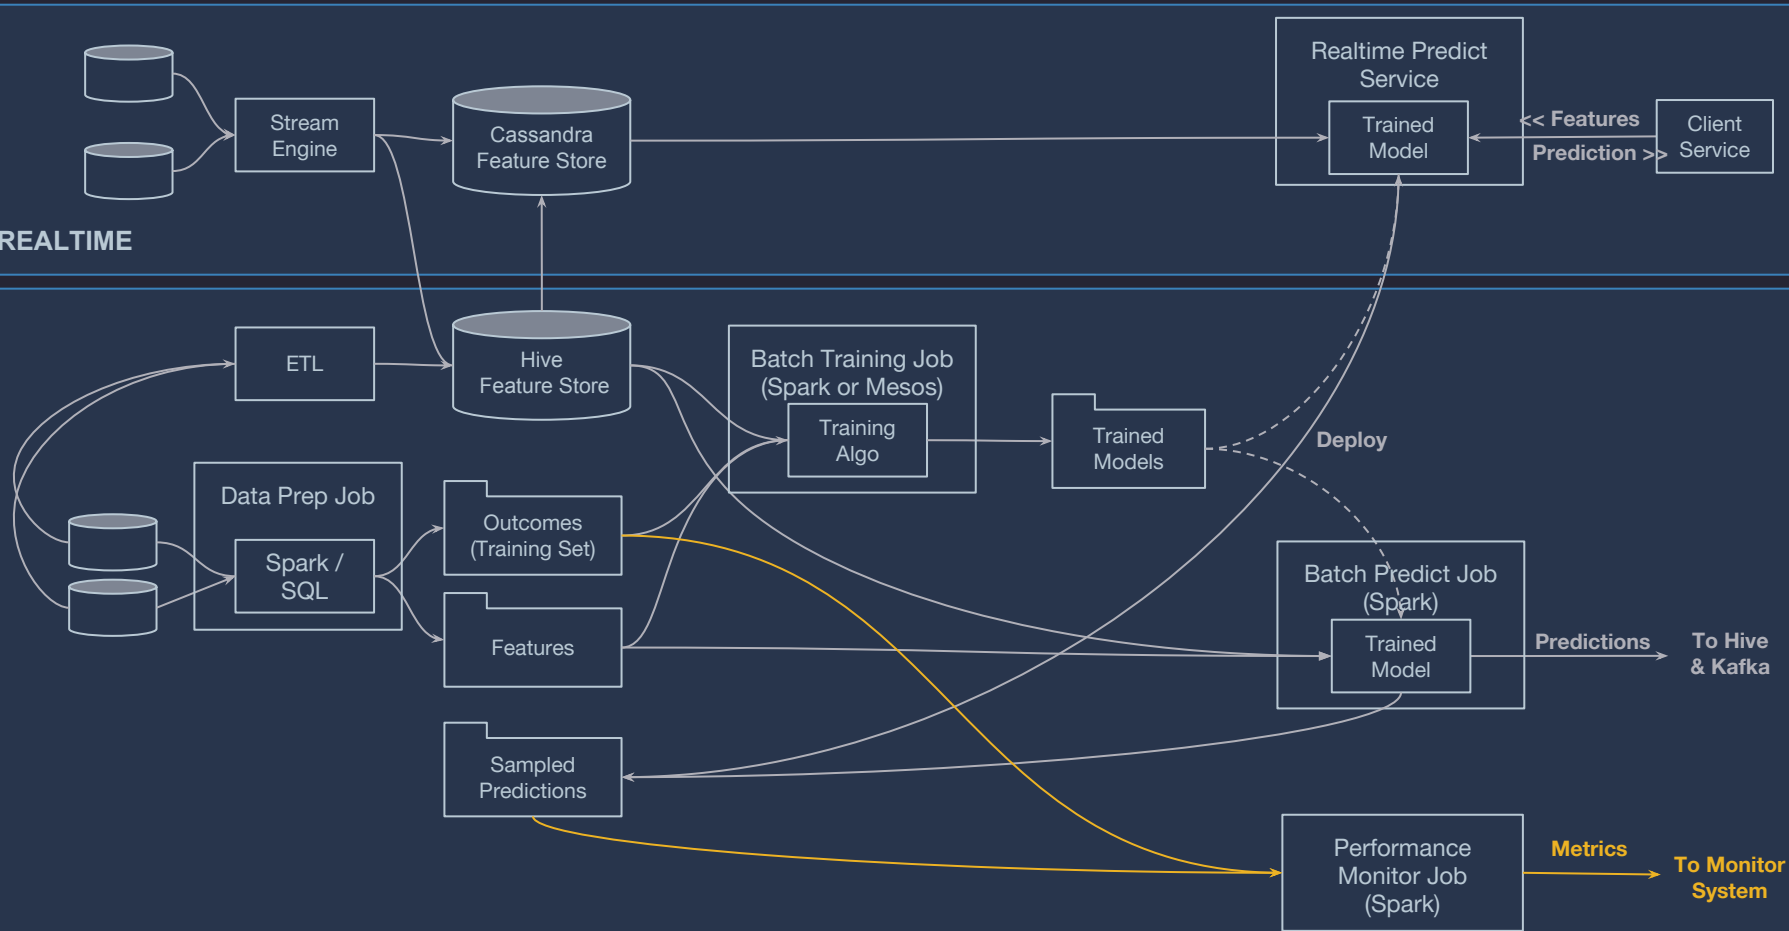
TRAIN MODELS

EVAL MODELS

DEPLOY, PREDICT &amp; MONITOR

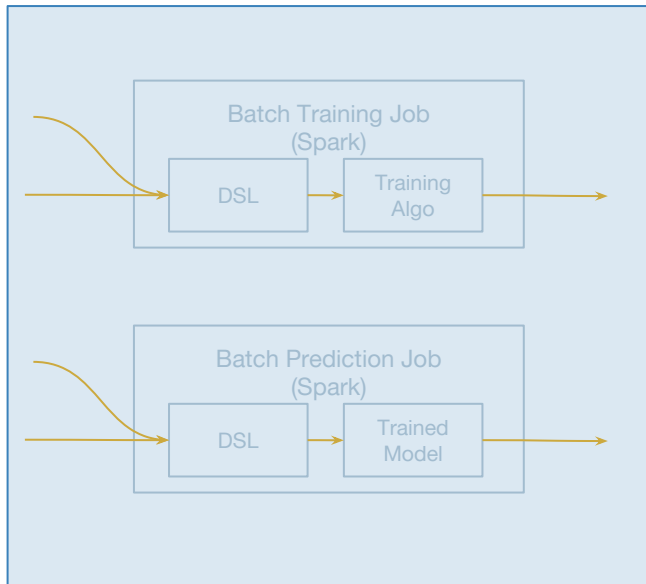
REALTIME

BATCH

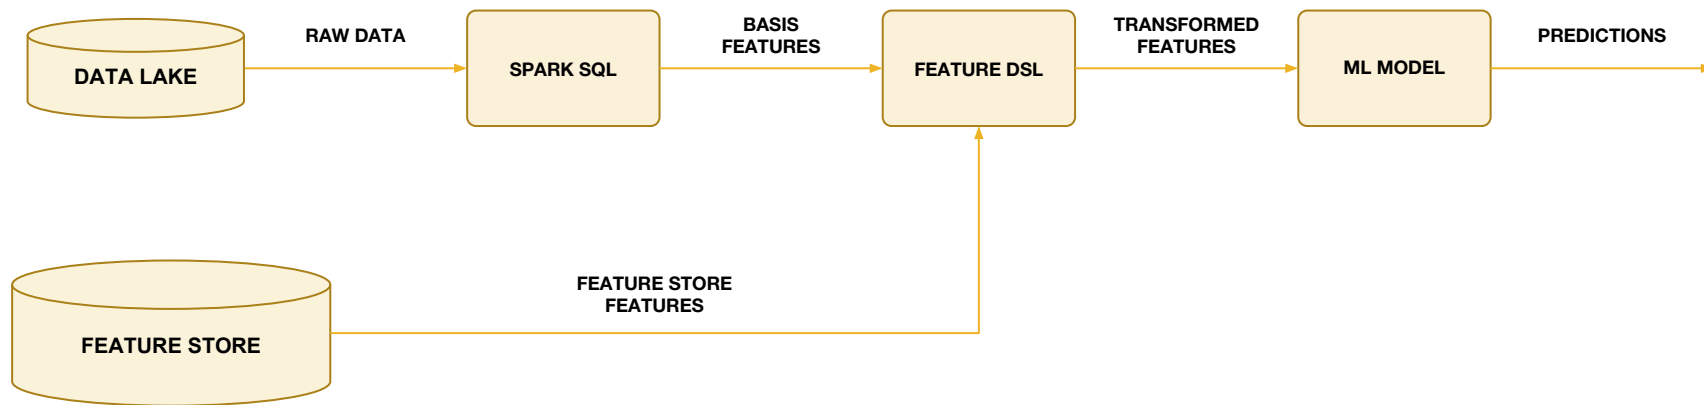


# DSL for Feature Engineering

- Pure function expressions for
  - Feature selection
  - Feature transformations (for derived and composite features)
- Standard set of accessor functions for
  - Feature store
  - Basis features
  - Column stats (min, max, mean, std-dev, etc)
- Standard transformation functions + UDFs
- Examples
  - @palette:store:orders:prep\_time\_avg\_1week:rs\_uuid
  - nFill(@basis:distance, mean(@basis:distance))



# Data Pipeline for Predictions



## Pain Points

- Users have to learn DSL for feature engineering
  - Waste of data scientist time
- ETL and Feature extraction could last for hours
  - Waste of data scientist time
- Features dump onto disk several times
  - Waste of data scientist time
- Feature store and data lake are separate
  - Waste of hardware resource
- Machine Learning and Data Analytics are separate
  - Waste of hardware resource and software resource

**Ideally**, Feature Engineering should ...

- Users do not need to learn anything new
- User do not need to worry about implementation details
- Machine Learning jobs are running as fast
- Features stored in Data Lake
- Features are cached in memory for repeated processing
- Machine Learning pipelines and Data Analytics share the same hardware, and **maybe software?**

## Try Declarative Language

- SQL extensible language
- Users no need to worry about implementation details
- Feature store resides in data lake
- Data kept in memory for repeated processing
- Machine Learning and Data Analytics
  - share the same hardware
  - share the same software
  - Streaming, Pipelined, Parallel, Vectorized execution



# The Data Model

- All datasets are tables
  - Training datasets
  - Validation datasets
- Labels and Features are columns
  - Label is integer
  - Feature is map<integer, float>
- Models are SQL extensible types
  - <size>:<model>
  - Model: Variable length serialized type
- Classify, Regress are SQL extensible functions

# The Execution Model

- Data are read from Hadoop once
- Data are kept in memory for repeated processing
- Streaming, Pipelined, Parallel, Vectorized execution
- Machine Learning Functions are implemented as:
  - Map
    - Applies to each input row
  - Combine
    - Merge partial aggregations
  - Reduce
    - Final aggregation

## In Summary

- Users use extensible SQLs to train and predict models
- All data are modeled as tables
- All data stored in data lake
- Data read from disk once, kept in memory for repeated processing
- A framework with execution primitives, to support:
  - SQL Analytics
  - Machine Learning


## What does it look like?

```
SELECT evaluate_classifier_predictions(label, classify(features, model))  
FROM (  
    SELECT learn_classifier(label, features) AS model  
    FROM training_data  
)  
CROSS JOIN validation_data
```

# The Presto Machine Learning Plugin

- All data are modeled as tables
- Data read from disk just once
  - Data are kept in memory for repeated processing
  - Streaming, Pipelined, Parallel, Vectorized execution
- Presto Connectors
  - SQL or Machine Learning on any data, e.g. Hadoop, MySQL, Elasticsearch, Cassandra, MongoDB, etc.
- Presto provides execution primitives:
  - Input Function -- **Map**
  - Combine Function -- **Combine**
  - Output function -- **Reduce**

# Significantly Reduce Complexity



- Data Scientists just use SQL for Machine Learning
- Model development time speed up > 5X
- Machine Learning pipeline speeds up > 10X
- Shared data lake
- Shared resource management
- Shared primitive executions

# What is Presto: Interactive SQL Engine for Big Data

Interactive query speeds

Horizontally scalable

ANSI SQL, Extensible for Machine Learning

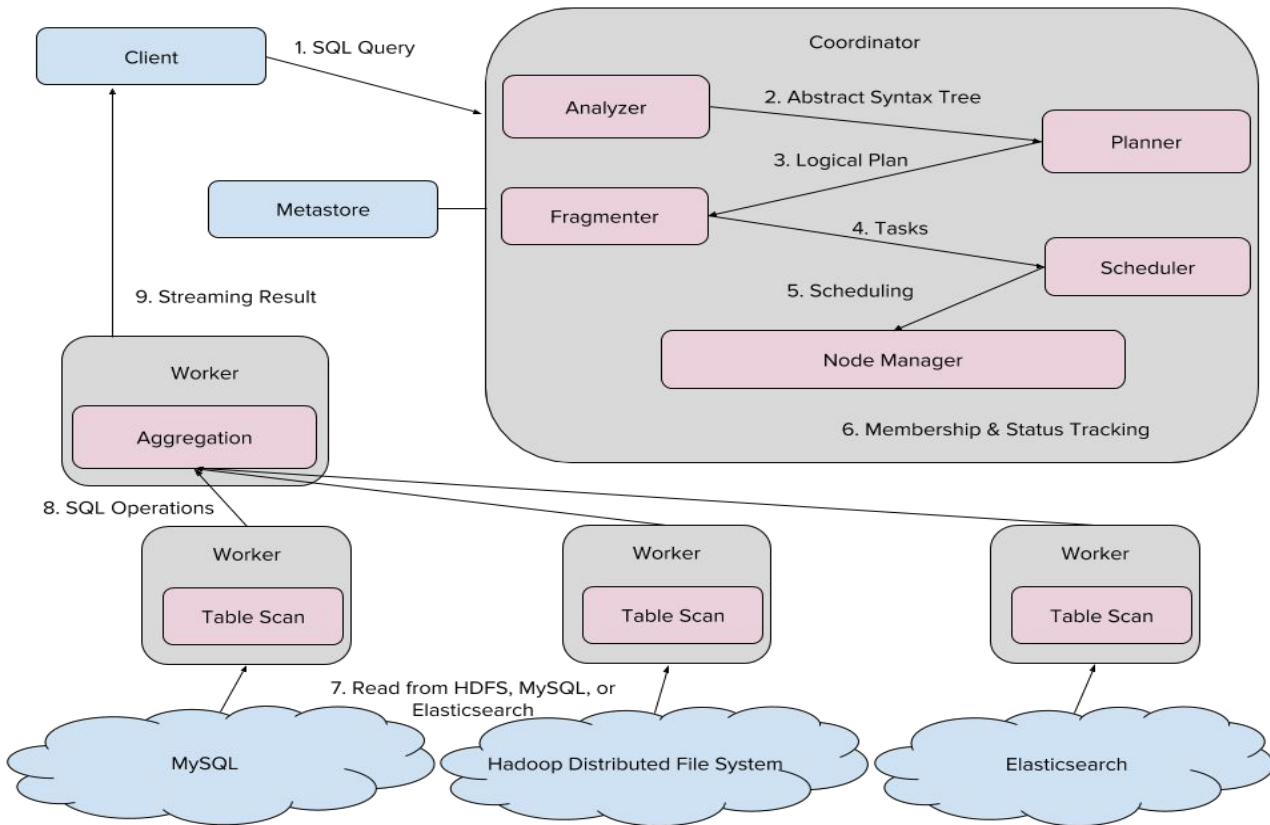
Battle-tested by Facebook, Uber, & Netflix

Completely open source

Access to petabytes of data in the Hadoop data lake



# How Presto Works



# Why Presto is Fast

- Data in memory during execution
- Pipelining and streaming
- Columnar storage & execution
- Bytecode generation
  - Inline virtual function calls
  - Inline constants
  - Rewrite inner loops
  - Rewrite type-specific branches

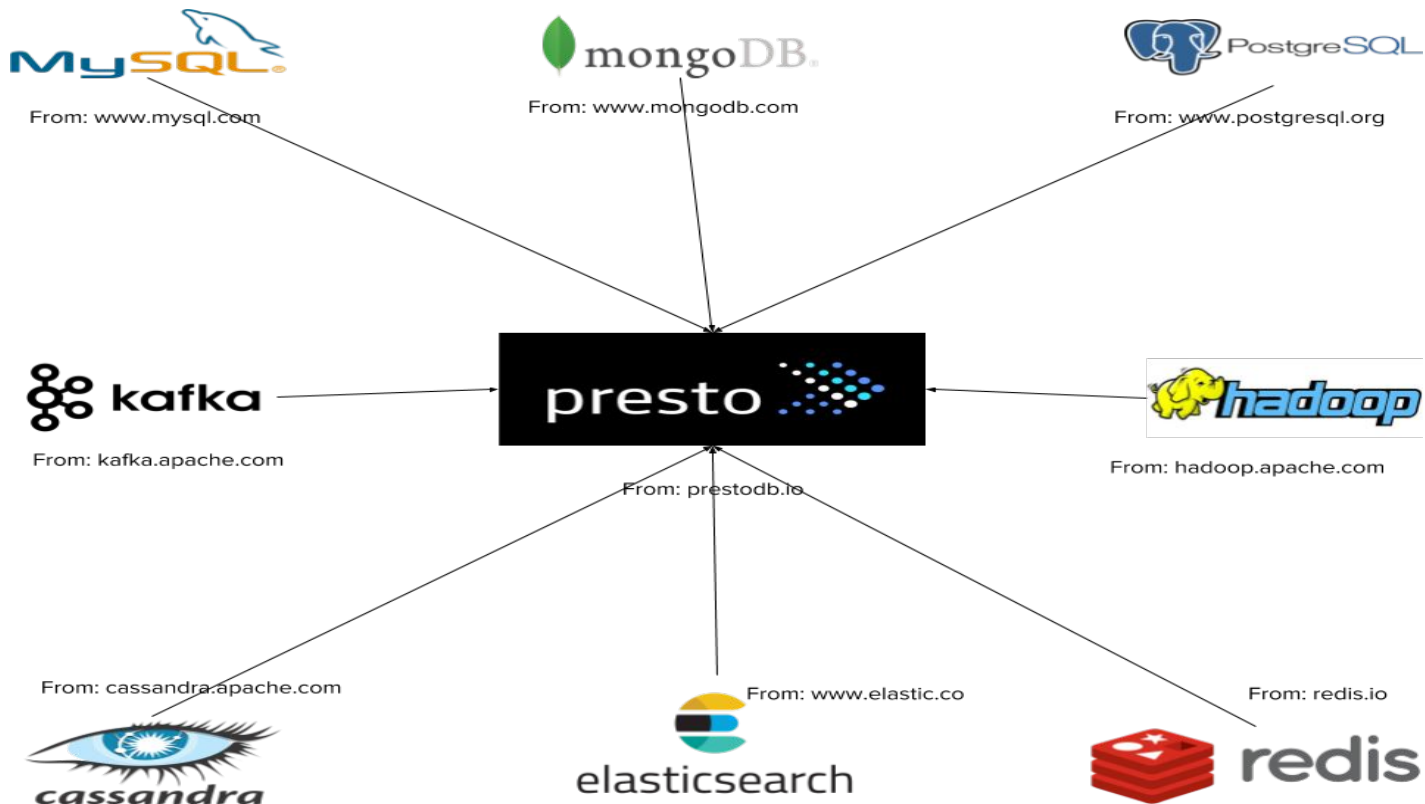
# Scale of Presto @ Uber

- 2 clusters
  - Application cluster
    - Hundreds of machines
    - 100K queries per day
    - **P90: 30s**
  - Ad hoc cluster
    - Hundreds of machines
    - 100K queries per day
    - **P90: 60s**
- Access to both raw and model tables
  - 5 petabytes of data
- Total **200K+** queries per day

# Resource Management

- Presto has its own resource manager
  - Not on YARN
  - Not on Mesos
- CPU Management
  - Priority queues
  - Short running queries higher priority
- Memory Management
  - Max memory per query per node
  - If query exceeds max memory limit, query fails
  - No OutOfMemory in Presto process

# Presto Connectors: No Need to Copy Data



# Presto Plugin

- Presto Connectors

- ConnectorMetadata
- ConnectorSplitManager
- ConnectorSplit
- ConnectorRecordCursor

- Extensible Types

- Native Container Type
- Type Signature

- Extensible Functions

- Input Function, Combine Function, Output Function

# Presto Machine Learning Plugin

- Types

- Model Type : variable length type
- Classifier Parametric Type

- Functions

- Classify
- Regress
- Features
- Learn\_libsvm\_regressor
- Learn\_libsvm\_classifier
- Evaluate\_classifier\_predictions



# Presto Ongoing Work

- Machine Learning support in Presto
- Presto Elasticsearch Connector
- Cost based optimizer
- Spill to disk
- Multi-tenancy Support
- Authentication and Authorization

We are Hiring  
<https://www.uber.com/careers/list/27366/>

Send resumes to:  
[luoz@uber.com](mailto:luoz@uber.com)

# Thank you

Interested in learning more about Uber Eng?  
Eng.uber.com  
Follow us on Twitter:  
@UberEng

Proprietary and confidential © 2016 Uber Technologies, Inc. All rights reserved. No part of this document may be reproduced or utilized in any form or by any means, electronic or mechanical, including photocopying, recording, or by any information storage or retrieval systems, without permission in writing from Uber. This document is intended only for the use of the individual or entity to whom it is addressed and contains information that is privileged, confidential or otherwise exempt from disclosure under applicable law. All recipients of this document are notified that the information contained herein includes proprietary and confidential information of Uber, and recipient may not make use of, disseminate, or in any way disclose this document or any of the enclosed information to any person other than employees of addressee to the extent necessary for consultations with authorized personnel of Uber.

The Uber logo, consisting of the word "UBER" in a bold, black, sans-serif font, is centered within a white square. The square is positioned in the bottom right corner of the slide, which has a teal background with a subtle geometric pattern of overlapping squares.